

Hui Xu

• (516) 457-4066 • huixucom@gmail.com • linkedin.com/in/huihsuxu • github.com/huixu11 • huixu11.github.io

EDUCATION

State University of New York at Stony Brook

MS in Engineering Artificial Intelligence — GPA: 3.67/4.0

Stony Brook, NY

Aug 2024 – Dec 2026

- **Coursework:** Distributed Systems, NLP, Deep Learning Algorithms, Robotics, Statistical Learning
- **Publication:** MAQuA: Adaptive Mental Health Screening using IRT (EACL 2026) [arXiv]

Beijing Forestry University

MS in Computer Application Technology (4.0/4.0) — BS in Information Systems (3.9/4.0)

Beijing, China

2009 – 2018

TECHNICAL SKILLS

Languages & Frameworks: Python, TypeScript, JavaScript, C, C++, Java, SQL, React, Redux, Django, Flask, FastAPI, PyTorch, TensorFlow, LangChain, PostgreSQL, SQLite, Redis, MySQL, AWS, GCP, Docker, Git, Jira, Ray, Multi-Paxos, Two-Phase Commit (2PC), DistAlgo, Consensus Algorithms

PROFESSIONAL EXPERIENCE

Mastercard

Beijing, China

Software Engineer II

Nov 2021 – Jul 2024

- Developed full-stack business analytics platform ("Test & Learn") for Chinese banks using **Django (backend)**, **React/Redux (frontend)**, and **PostgreSQL/SQLite**, enabling local deployment and data compliance
- Proposed and prototyped migration from in-house multiprocessing framework to **Ray Core**, enabling **distributed execution across clusters** and containerized environments with minimal code changes
- Engineered high-performance **outlier detection algorithm** using statsmodels leave-one-out statistics; **reduced runtime from 9 hours to 10 minutes (98.1% improvement)** through selective computation and vectorization
- Built **Metric Uploader** feature capable of processing **400MB+ CSV files in under 60 seconds** with row-level validation using fully vectorized algorithms and comprehensive unit testing coverage
- Designed and delivered Driver Summary module featuring driver significance analysis and visual summaries using React, Redux hashmaps, and Recharts; collaborated with Product Managers and Tech Leads to refine requirements

Dazhangfang (Chinese Intuit)

Beijing, China

Python Engineer

Jul 2018 – Oct 2021

- Managed and deployed large-scale **OCR platform (100,000+ lines of code, 10 servers)** integrating recognition engine, invoice verification service, and web service for automated receipt recognition and accounting
- **Optimized database queries and indexing**, improving Invoice Recognition Web Service performance by **99.99%**, dramatically reducing response latency
- Automated invoice verification process, achieving **90% reduction in human intervention** using edit-distance algorithms for text matching and validation
- Implemented **asynchronous task scheduling and message delivery** using APScheduler and Redis as message queue broker, increasing throughput and system reliability

PROJECTS

AgentBeats – Financial AI Benchmark & Agent Security | Team Lead, Python, FastAPI, LLMs

- **Leading 5-member team** building dual-agent adversarial evaluation system for AI financial analysts using Python, FastAPI, LLMs, Multi-Agent Systems, MCP, and LangChain
- Integrating BizFinBench.v2 (29K+ Q&A pairs) and designing **Alpha Score metric** for reasoning robustness evaluation
- Developing **FastAPI services and MCP servers** for real-time SEC and Yahoo Finance data access; analyzing prompt injection and LLM tool-use vulnerabilities (LangChain CVE) for agent security research

Scalable Distributed Transaction System | DistAlgo, SQLite, Multi-Paxos, 2PC

- Designed and implemented **9-node, 3-cluster replicated transaction system** with Multi-Paxos for intra-shard consensus and **Two-Phase Commit (2PC)** protocol for cross-shard atomic transactions
- Achieved **sub-1.5s leader failover** through heartbeat-based failure detection and automatic re-election; implemented configurable cluster topology (N clusters × M nodes) with automatic shard map generation

Transformer Language Model from Scratch | Python, PyTorch, NLP, Transformers

- Engineered Transformer LM from scratch with **BPE tokenizer, RMSNorm, RoPE, multi-head attention, and SwiGLU layers**; built and tested Linear, Embedding, and Attention modules ensuring stable gradients
- Trained on TinyStories/OpenWebText using multiprocessing pre-tokenization and custom **AdamW optimizer**; achieved sub-2 minute BPE training (10K vocab) with fluent text generation and competitive perplexity

ACHIEVEMENTS

Jane Street GPU Mode Hackathon: 10th Place - Optimized latency and accuracy using dynamic batching strategies

Hack-Nation Hackathon: ThoughtLink 1st Place: decoding high-level human intent from non-invasive brain signals and mapping it directly to executable robot instructions.

Agentic AI MOOC fall 2025 legendary tier UC Berkeley, Center for Responsible, Decentralized Intelligence