

HUI XU

Stony Brook, NY, United States of America

📞 (516)457-4066 📩 huixucom@gmail.com 💬 linkedin.com/in/huihsuxu/

🌐 huixu11.github.io

EDUCATION

Masters in Engineering Artificial Intelligence - GPA: 3.67/4	Aug 2024 - Dec 2025
<i>State University of New York - Stony Brook</i>	<i>NY, USA</i>
Coursework: Distributed System, Natural Language Processing , Deep Learning Algorithms and Software	
MAQuA: Adaptive Question-Asking for Multidimensional Mental Health Screening using Item Response Theory	
Accepted by EACL (view paper)	
Masters in Computer Application Technology - GPA: 4/4	Sep 2013 - June 2018
<i>Beijing Forestry University</i>	<i>Beijing, China</i>
Bachelor in Information Management and Information Systems - GPA: 3.9/4	Sep 2009 - June 2013
<i>Beijing Forestry University</i>	<i>Beijing, China</i>

WORK EXPERIENCE

Mastercard Django, React, Python, Redux, Typescript, PostgreSQL, SQLite, Ray	Nov 2021 - Jul 2024
<i>Software Engineer II</i>	<i>Beijing, China</i>
<ul style="list-style-type: none">Built a full-stack business analytics platform (“Test & Learn”) from the ground up for Chinese banks, enabling local deployment and data compliance, using Django (backend), React/Redux (frontend), and PostgreSQL/SQLite.Proposed and led migration from an in-house multiprocessing framework to Ray Core, enabling distributed execution across clusters and containerized environments with minimal code changes.Developed a high-performance outlier detection algorithm using statsmodels leave-one-out statistics; improved runtime from 9 hours to 10 minutes through selective computation and vectorization.Designed and implemented a script generator for large-scale dummy datasets with O(1) space complexity, significantly improving data-simulation efficiency for testing.Led the design and delivery of a Driver Summary module showing driver significance and visual summaries via React, Redux hashmaps, and Recharts; collaborated with PMs and tech leads to refine product requirements and architecture.Architected and implemented a Metric Uploader feature capable of processing 400MB+ CSV files in under one minute, with row-level validation using a fully vectorized algorithm and comprehensive unit testing coverage.Designed data structures and algorithms for dynamic result grid manipulation using react-beautiful-dnd and Redux, enabling flexible drag-and-drop analytics views.Demonstrated strong leadership and ownership—regularly identified technical roadblocks, coordinated across engineering and product teams, and drove feature completion under tight timelines.	
Dazhangfang (Chinese Intuit) SQL, Redis, APScheduler, OCR, Flask, Google Cloud July 2018 - Oct 2021	
<i>Python Engineer</i> <i>Beijing, China</i>	
<ul style="list-style-type: none">Maintained and enhanced the company’s invoice and bank form recognition system and finance/taxation APIs, enabling users to upload receipts for automatic recognition, classification, and accounting.Managed and deployed a large-scale recognition platform (100,000 lines of code, 10 servers) integrating the recognition engine, invoice verification service, and web service for recognized results.Optimized database queries and indexing, improving the Invoice Recognition Web Service performance by 99.99%, dramatically reducing response latency.Automated manual invoice verification, achieving a 90% reduction in human intervention using edit-distance algorithms for text matching and validation.Implemented asynchronous task scheduling and message delivery using APScheduler and Redis as a message queue broker, increasing throughput and reliability.Designed caching mechanisms for recognition results and optimized the end-to-end OCR pipeline, significantly reducing compute cost and improving system scalability.	

PROJECTS

Transformer Language Model from Scratch (CS336) | Python, PyTorch, NLP, BPE, Transformer, CUDA

- <https://github.com/helperfunc/assignment1-basics>
- Engineered a Transformer LM from scratch with BPE tokenizer, RMSNorm, RoPE, multi-head attention, and SwiGLU layers.
- Built and tested Linear, Embedding, and Attention modules ensuring stable gradients and efficient backprop.
- Trained on TinyStories/OpenWebText using multiprocessing pre-tokenization and custom AdamW optimizer.
- Achieved < 2 min BPE training (10 K vocab) and fluent text generation with competitive perplexity.

Jane Street GPU Mode Hackathon 10th Place | Python, PyTorch

- <https://github.com/helperfunc/hackathon>
- Optimized latency and accuracy of the code using **dynamic batching strategies**.
- Achieved About 71.1ms latency and 0.77 accuracy, earn \$127/s.